




Relationship between one processor and inputs, compared to speedup




Relationship between two processors and inputs, compared to speedup




Relationship between three processors and inputs,
compared to speedup



Relationship between four processors and inputs, compared
to speedup



Relationship between the number of processors and inputs,
compared to speedup



Results by processor count (analysing speedup)

1 process (p=1).

Speedup is < 1 across all input sizes ($\sim 0.84\text{--}0.94$). With no parallelism, PSRS still pays its fixed costs (sampling, pivot handling, partition bookkeeping), so it is slower than straight quicksort.

2 processes (p=2).

Speedup improves to $\sim 1.62\text{--}1.70$ as N grows. Once each rank has enough local work, the cost of one all-to-all exchange and the final merge is properly distributed. This is the first routine where PSRS consistently beats sequential quicksort.

3 processes (p=3).

Speedup rises to $\sim 2.04\text{--}2.40$. Communication volume grows with p , but for $N \geq 10^5$ the extra concurrency outweighs it. Curves begin to bend, as because pivot selection, all-to-all, and local merge phases do not scale perfectly.

4 processes (p=4).

Best overall times: $\sim 2.50\text{--}3.15\times$ speedup at $N=10^6\text{--}10^7$. Scaling is good but not linear (ideal would be $4\times$). Diminishing returns come from:

- 1) More messages in the all-to-all
- 2) Imperfect load balance from regular sampling
- 3) The final k-way merge, which is $O(n \log p)$ locally.

Interpretation for how PSRS steps affect speedup

- **Step 1 (local sort).** Helpful to scale: each rank sorts $\sim N/p$ elements in $O((N/p) \log(N/p))$. This is where most of the parallel savings come from.
- **Step 2–3 (sampling + pivoting).** Overhead grows with p (rank 0 sorts p^2 samples and broadcasts $p-1$ pivots). For small N this dominates; for large N it's negligible but still caps speedup as p increases.
- **Step 4 (partition + all-to-all).** Dominant communication cost. Each rank sends/receives $O(N/p)$ data split across p peers. Latency (many small messages) and bandwidth (total bytes) both increase with p .

- **Step 5 (k-way merge).** Each rank merges p sorted runs totaling $\sim N/p$. Even with an $O(n \log p)$ merge, this is a serial section on each process and contributes to the non-linear scaling at higher p .

Takeaways

1. PSRS is slower than sequential quicksort for small N due to fixed overheads (sampling, pivoting, all-to-all, merge).
2. For moderate–large N , PSRS outperforms sequential quicksort and strong scaling up to 4 processes ($\approx 3.15x$ speedup at 10^7 inputs).
3. Sub-linear scaling is expected: Amdahl's serial fractions (pivoting/merge) and the communication pattern (All to all) limit ideal speedup. Load imbalance from regular sampling further reduces efficiency.

This aligns with the PSRS design goal for MIMD systems: good performance for large inputs with a simple, analyzable communication pattern, and predictable diminishing returns as p increases.